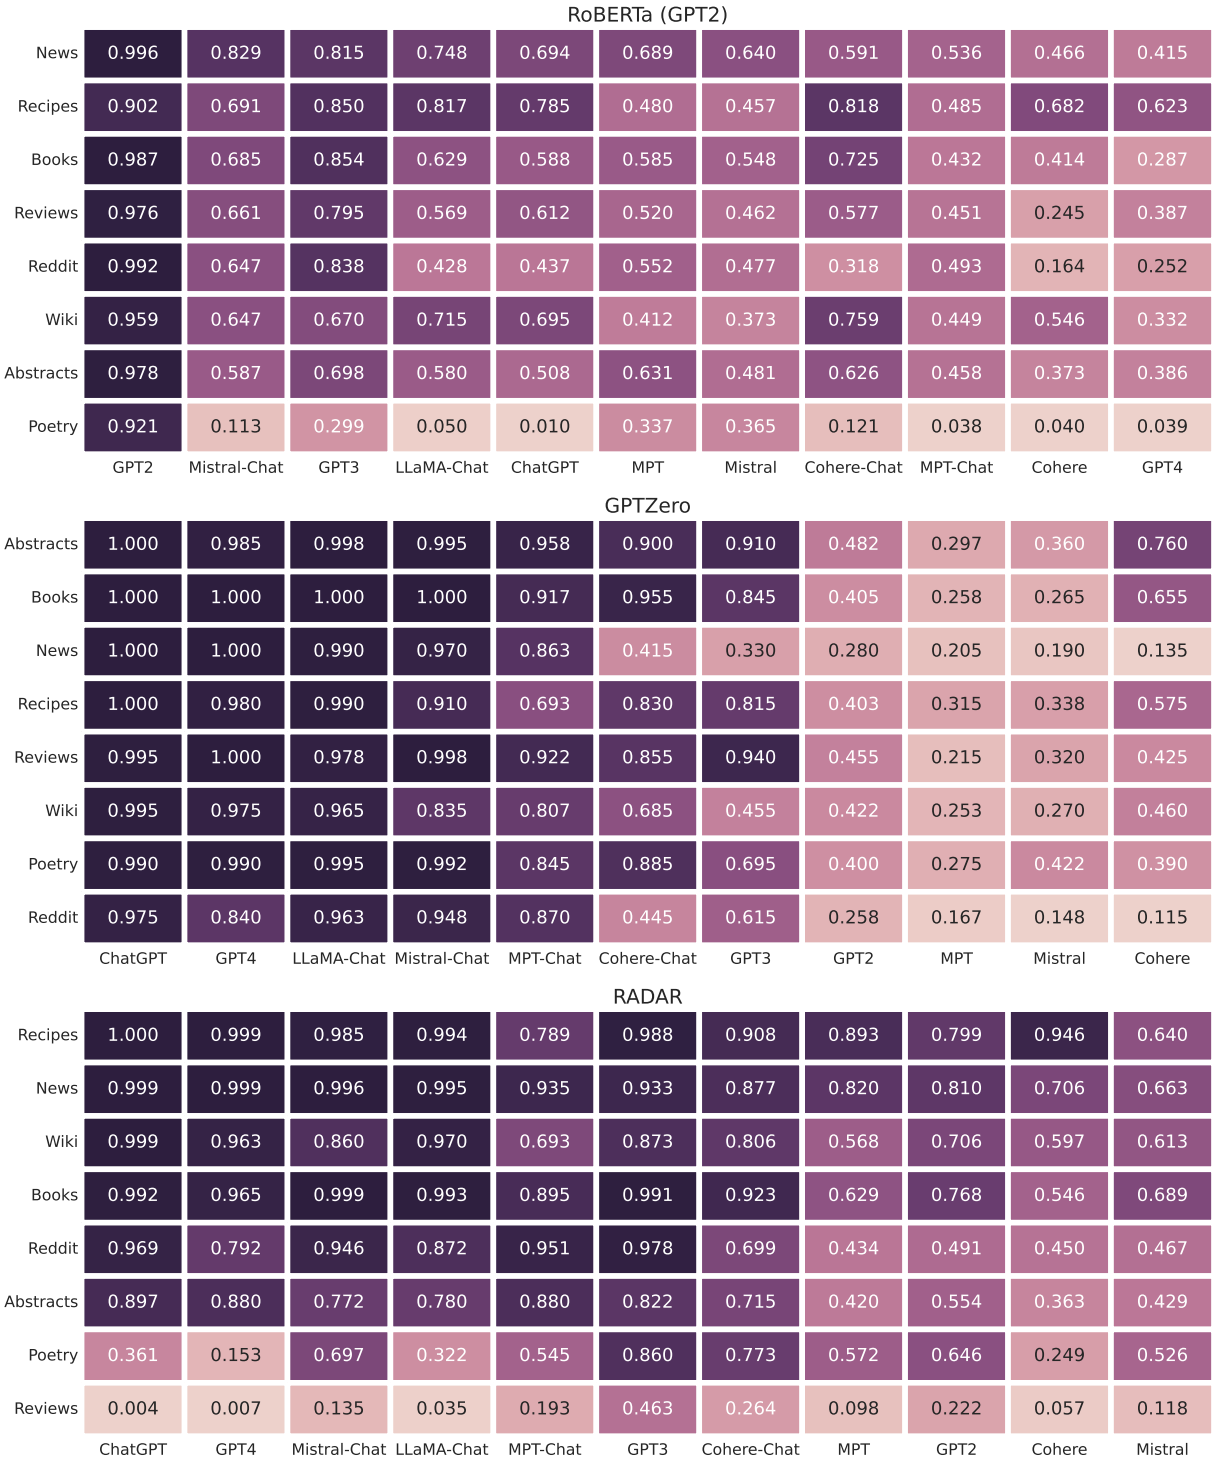

Figure 12: Extended heatmap of RoBERTa GPT2, GPTZero, and RADAR's performance across all models and domains in the RAID dataset. We see that the trends noted in Figure 6 still hold.

	News	Wiki	Reddit	Books	Abstracts	Reviews	Poetry	Recipes
RoBERTa-B GPT2	0.032 (5.0%)	0.379 (5.0%)	0.477 (5.0%)	0.586 (5.0%)	0.055 (5.0%)	0.539 (5.0%)	0.998 (5.0%)	0.916 (5.0%)
RoBERTa-L GPT2	0.070 (5.0%)	0.459 (5.0%)	0.100 (5.0%)	0.161 (5.0%)	0.085 (5.0%)	0.298 (5.1%)	0.762 (5.0%)	0.315 (5.0%)
RoBERTa-B C-GPT	0.987 (5.0%)	0.983 (5.0%)	0.219 (5.0%)	0.996 (5.0%)	0.007 (5.0%)	0.371 (5.0%)	0.295 (5.0%)	0.998 (5.0%)
RADAR	0.022 (5.0%)	0.061 (5.0%)	0.695 (5.0%)	0.174 (5.0%)	0.31 (5.0%)	0.997 (5.0%)	0.457 (5.0%)	0.016 (5.0%)
GLTR	0.788 (5.0%)	0.788 (5.0%)	0.767 (5.0%)	0.742 (5.0%)	0.726 (5.0%)	0.757 (5.0%)	0.756 (5.0%)	0.863 (5.0%)
FastDetectGPT	0.920 (4.8%)	0.870 (5.1%)	0.870 (5.1%)	0.930 (4.9%)	0.860 (4.6%)	0.900 (4.8%)	0.940 (5.9%)	0.880 (5.5%)
LLMDet	1.000 (5.0%)	1.000 (5.0%)	1.000 (5.0%)	1.000 (5.0%)	0.999 (5.0%)	1.000 (5.1%)	1.000 (5.0%)	0.998 (5.0%)
Binoculars	0.077 (4.9%)	0.093 (5.0%)	0.099 (5.0%)	0.085 (5.0%)	0.092 (5.0%)	0.097 (4.9%)	0.084 (5.0%)	0.094 (5.0%)
GPTZero	0.047 (5.0%)	0.032 (5.0%)	0.057 (5.0%)	0.125 (5.0%)	0.125 (5.0%)	0.070 (5.0%)	0.031 (5.0%)	0.035 (5.0%)
Originality	0.375 (5.0%)	0.938 (5.0%)	0.250 (5.0%)	0.312 (5.0%)	0.257 (5.0%)	0.461 (5.0%)	0.047 (5.0%)	0.750 (5.0%)
Winston	0.001 (4.0%)	0.970 (5.0%)	0.062 (5.0%)	0.998 (5.0%)	0.000 (6.0%)	0.062 (5.0%)	0.875 (5.0%)	0.996 (5.0%)
ZeroGPT(*)	1.000 (29.0%)	1.000 (48.0%)	0.375 (1.0%)	0.250 (9.0%)	0.500 (4.0%)	1.000 (5.0%)	0.125 (5.0%)	1.000 (52.0%)

Table 13: Thresholds found by our search and the exact False Positive Rates on our dataset. We see that ZeroGPT is incapable of achieving the target FPR of 5% in many domains.

	News	Wiki	Reddit	Books	Abstracts	Reviews	Poetry	Recipes
RoBERTa-B GPT2	74.3	64.7	56.3	67.9	56.3	69.1	23.9	64.6
RoBERTa-L GPT2	69.8	59.5	54.1	62.4	58.9	58.2	24.4	67.2
RoBERTa-B CGPT	45.1	48.7	44.6	53.5	72.3	65.3	32.0	5.9
RADAR	88.0	76.8	71.8	84.5	66.7	14.1	53.0	**88.5**
GLTR	66.9	64.3	65.7	74.0	62.0	67.3	34.8	67.2
FastDetectGPT	74.2	77.3	70.9	76.2	76.4	77.5	63.4	74.6
LLMDet	39.8	32.6	39.6	37.1	18.0	33.1	30.7	48.1
Binoculars	80.7	76.7	79.4	83.7	79.1	80.1	**81.0**	76.6
GPTZero	58.1	62.8	57.0	71.4	74.9	70.5	69.5	67.6
Originality	**88.4**	**83.2**	**85.0**	**90.4**	87.7	**87.3**	75.1	82.8
Winston	72.4	54.9	68.9	70.7	**94.7**	72.9	64.3	68.9
ZeroGPT(*)	72.2	70.6	65.1	73.3	60.3	68.6	50.0	63.7

Table 14: Accuracy Score at FPR=5% for detectors across different domains. We see that metric-based methods perform surprisingly well across domains and that detectors can perform surprisingly poorly on unseen domains.

	GPT2	GPT3	ChatGPT	GPT4	Cohere		Mistral		MPT		Llama	Total
Chat? (Y/N)	✗	✗	✓	✓	✗	✓	✗	✓	✗	✓	✓	-
R-B GPT2	84.0	74.7	65.4	42.4	42.9	61.1	45.7	65.9	49.2	45.9	68.7	59.1
R-L GPT2	96.3	72.4	53.7	33.8	37.3	56.6	47.6	60.5	52.6	41.6	56.7	56.7
R-B CGPT	36.7	54.2	66.1	30.0	31.3	49.6	17.7	69.9	17.8	53.3	70.1	44.8
RADAR	64.7	88.6	82.1	76.0	51.4	77.2	54.1	83.6	58.0	76.5	78.5	70.9
GLTR	66.6	85.1	81.4	53.7	54.2	67.4	49.1	75.4	35.4	52.5	81.6	62.6
F-DetectGPT	72.1	95.4	96.1	73.9	84.7	85.1	58.2	81.3	45.3	57.1	94.0	73.6
LLMDet	48.2	40.2	18.9	27.0	32.6	35.6	31.5	35.7	28.2	21.4	55.1	35.0
Binoculars	68.9	**99.2**	**99.6**	91.9	**94.8**	**95.4**	62.3	91.7	45.2	70.8	97.6	79.6
GPTZero	38.8	70.1	99.4	97.1	43.9	74.6	28.9	**95.6**	24.8	85.9	**98.5**	66.5
Originality	**99.1**	98.2	98.2	89.9	78.9	90.6	**71.0**	95.5	**58.1**	76.2	94.7	**85.0**
Winston	47.6	77.8	**99.6**	**98.8**	63.6	86.2	46.1	94.9	24.8	**79.2**	97.5	71.0
ZeroGPT(*)	42.4	90.2	93.2	67.1	65.9	76.6	49.3	81.4	27.3	66.0	93.7	65.5

Table 15: Accuracy at FPR=5% for detectors on non-adversarial outputs of different models. We see that base models are more difficult to detect than their chat fine-tuned counterparts and that metric-based methods show impressive cross-model generalization. Asterisks (*) indicate that the detector was unable to achieve the target FPR.

	None	AS	AD	HG	IP	NS	PP	MS	SYN	ULS	WSA	ZWS
RoB-B GPT2	59.1	55.6	37.1	7.6	56.9	55.9	68.9	43.8	71.5	18.8	45.2	99.9
RoB-L GPT2	56.7	52.4	33.2	21.3	55.1	51.7	72.9	39.5	79.4	19.3	40.1	99.9
RoB-B CGPT	44.8	43.3	38.0	0.0	5.2	44.3	49.2	42.1	39.6	31.7	0.1	0.0
RADAR	70.9	70.8	67.9	59.3	73.7	71.0	67.3	69.5	67.5	70.4	66.1	82.2
GLTR	62.6	61.2	52.1	24.3	61.4	59.9	47.2	59.8	31.2	48.1	45.8	97.2
F-DGPT	73.6	71.6	64.7	51.4	72.0	68.2	71.8	70.7	34.0	60.4	64.4	98.9
LLMDet	35.0	33.9	27.4	40.6	27.2	33.8	28.5	32.7	27.3	23.4	4.4	27.1
Binoculars	79.6	78.2	74.3	37.7	71.7	77.1	80.3	78.0	43.5	73.8	70.1	99.1
GPTZero	66.5	64.9	61.0	66.2	66.2	65.8	64.0	65.1	61.0	56.5	66.2	66.2
ZeroGPT	65.5	65.4	59.7	82.4	64.9	64.7	46.7	64.7	18.8	54.5	64.2	48.0
Originality	85.0	83.6	71.4	9.3	85.1	86.0	96.7	78.6	96.5	75.8	84.9	4.9
Winston	71.0	68.9	66.9	26.3	69.8	69.0	52.6	67.5	63.6	56.8	46.8	25.0

Table 16: Accuracy Score at FPR=5% for all detectors across different adversarial attacks. Abbrevations are: AS: Alternative Spelling, AD: Article Deletion, HG: Homoglyph, IP: Insert Paragraphs, NS: Number Swap, PP: Paraphrase, MS: Misspelling, SYN: Synonym Swap, ULS: Upper Lower Swap, WSA: Whitespace Addition, ZWS: Zero-Width Space Addition